

```
from google.colab import files
import io
import pandas as pd
data = files.upload()

Choose Files intern_data_spm.csv
• intern_data_spm.csv(text/csv) - 1046715 bytes, last modified: 6/6/2023 - 100% done
Saving intern_data_spm.csv to intern_data_spm (1).csv
```

```
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from datetime import datetime
%matplotlib inline
```

```
df=pd.read_csv(io.StringIO(data['intern_data_spm.csv'].decode('utf-8')))
```

```
df.shape
```

(5000, 16)

```
df.head()
```

	CLOCK	X1	X2	X3	X4	X5	X6	X7	X8	X9	
0	2022-12-18 14:50:00	48.65105	38.97555	35.33245	36.00630	63.58980	10.96650	145.15495	384.60800	919.95	2
1	2022-12-31	46.30760	36.69085	32.61530	34.56650	62.49190	10.48560	153.96970	389.11030	880.05	1
2	2022-12-23	48.56460	38.86065	34.98840	35.50790	65.73585	10.72395	161.72370	383.47905	712.95	1

```
df.tail()
```

	CLOCK	X1	X2	X3	X4	X5	X6	X7	X8	X9	
4995	2022-12-29 17:50:00	48.970250	38.739500	35.964150	36.376850	65.717900	8.880750	152.399100	408.22870		
4996	2022-12-29 00:40:00	47.482000	38.186750	34.010450	35.813450	64.686100	17.171500	134.963450	368.41175		
4997	2022-12-19	47.177000	38.093600	34.074500	35.629350	63.249700	11.424150	135.836050	378.53305		

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CLOCK      5000 non-null   object
1   X1          5000 non-null   float64
2   X2          5000 non-null   float64
3   X3          5000 non-null   float64
4   X4          5000 non-null   float64
5   X5          5000 non-null   float64
6   X6          5000 non-null   float64
7   X7          5000 non-null   float64
8   X8          5000 non-null   float64
9   X9          5000 non-null   float64
10  X10         5000 non-null   float64
11  X11         5000 non-null   float64
12  X12         5000 non-null   float64
13  X13         5000 non-null   float64
14  X14         5000 non-null   float64
15  output      5000 non-null   float64
dtypes: float64(15), object(1)
memory usage: 625.1+ KB
```

```
df.isnull().sum()
```

CLOCK	0
X1	0
X2	0
X3	0

```

X4      0
X5      0
X6      0
X7      0
X8      0
X9      0
X10     0
X11     0
X12     0
X13     0
X14     0
output  0
dtype: int64

```

```
[features for features in df.columns if df[features].isnull().sum()>0]
```

```
[]
```

NO NULL VALUES SO THERE ARE NO MISSING VALUES

```
df.dtypes
```

```

CLOCK    object
X1      float64
X2      float64
X3      float64
X4      float64
X5      float64
X6      float64
X7      float64
X8      float64
X9      float64
X10     float64
X11     float64
X12     float64
X13     float64
X14     float64
output   float64
dtype: object

```

Saving...




```

output   float64
dtype: object

```

```
df.plot()
```


<Axes: >



```
df.sort_values(['CLOCK', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X7', 'X8', 'X9', 'X10', 'X11', 'X12', 'X13', 'X14', 'output'])
```

	CLOCK	X1	X2	X3	X4	X5	X6	X7	X8	X9
1703	2022-12-10 06:30:00	49.52130	39.18235	35.92655	36.67485	63.90860	2.63820	142.41390	387.24310	920.00
3094	2022-12-10 06:50:00	49.55760	39.07140	35.79370	36.81315	62.28685	2.79050	146.99295	402.32190	920.05

```
df['output'].plot(figsize=(12,4))
```



```
cols=df.columns
df[['X3', 'X4', 'X5', 'X6', 'X7', 'X8', 'X9', 'X10', 'X11', 'X12', 'X13', 'X14', 'output']].duplicated()
duplicated_rows=df[duplicates]
print(duplicated_rows)
print(duplicates)
```

	CLOCK	X1	X2	X3	X4	X5	X6	X7	X8	X9	X10	X11	X12	X13	X14	output
52	2022-12-13 06:00:00	48.111000	39.059250	34.437700	37.116600	62.575800	17.052400	132.687400	358.084200	919.950000	2.17830	2.829300	3.804500	3.799000	7.694000	38.972250
107	2022-12-17 01:20:00	46.827900	37.523900	33.374050	35.422150	65.338500	10.199800	136.478150	369.819700	919.950000	1.93795	2.679350	2.913000	2.923500	7.693500	38.436000
108	2022-12-19 17:00:00	47.433500	38.433700	34.218200	35.896800	64.227100	8.185350	136.866000	363.454600	920.000000	2.05755	2.783350	3.659000	3.665000	7.693300	39.025950
118	2022-12-15 10:00:00	47.569650	38.379900	34.173000	36.184050	66.840300	8.052250	138.086150	379.072150	919.950000	1.80095	2.540200	3.097000	3.095500	7.693850	36.270600
145	2022-12-13 06:00:00	48.111000	39.059250	34.437700	37.116600	62.575800	17.052400	132.687400	358.084200	919.950000	2.17830	2.829300	3.804500	3.799000	7.694000	38.972250
...
4994	2022-12-25 07:40:00	46.379550	36.759050	33.223450	34.295950	64.590650	6.571300	156.399550	413.743950	920.000000	1.68955	2.508800	2.712500	2.703000	7.693900	38.868850
4995	2022-12-29 17:50:00	48.970250	38.739500	35.964150	36.376850	65.717900	8.880750	152.399100	408.228700	909.900000	1.79045	2.555800	2.760500	2.788500	7.693900	37.104800
4996	2022-12-29 00:40:00	47.482000	38.186750	34.010450	35.813450	64.686100	17.171500	134.963450	368.411750	919.950000	1.89235	2.610600	2.799000	2.839500	7.693900	38.265500
4997	2022-12-19 03:20:00	47.177000	38.093600	34.074500	35.629350	63.249700	11.424150	135.836050	378.533050	920.050000	2.04950	2.748300	3.621500	3.631500	7.693300	38.432150
4998	2022-12-12 13:50:00	48.468167	39.066556	34.451111	36.946222	62.825278	10.619111	135.960722	368.527111	919.944444	2.13950	2.801833	3.478889	3.467222	7.693833	35.841333

[2539 rows x 16 columns]


0	False
1	False

```
2      False
3      False
4      False
...
4995   True
4996   True
4997   True
4998   True
4999   False
Length: 5000, dtype: bool
```

CHEKING DUPLICATES

```
duplicates = df.duplicated()
df1 = df[duplicates == True]
df2 = df[duplicates == False]
x = ['Duplicates', 'Non-Duplicates']
y = [len(df1), len(df2)]

# Create a bar plot
plt.figure(figsize=(8, 6))
plt.bar(x, y)
plt.xlabel('Duplicate Status')
plt.ylabel('Count')
plt.title('Duplicates in DataFrame')
plt.show()
```



```
#REMOVING DUPLICATES
df=df.drop_duplicates()
print(df)
```

	CLOCK	X1	X2	X3	X4	X5	\
0	2022-12-18 14:50:00	48.65105	38.97555	35.33245	36.00630	63.58980	
1	2022-12-31 12:10:00	46.30760	36.69085	32.61530	34.56650	62.49190	
2	2022-12-23 18:10:00	48.56460	38.86065	34.98840	35.50790	65.73585	
3	2022-12-19 05:30:00	47.59175	38.35315	34.27105	35.68830	64.44080	
4	2022-12-10 14:00:00	49.20060	38.82830	35.15390	36.69245	63.76605	
...	
4987	2022-12-21 17:10:00	46.74265	37.91030	33.77445	35.46620	62.01295	
4988	2022-12-20 13:30:00	46.81005	37.49205	34.34830	35.62205	62.44480	
4989	2022-12-24 04:50:00	46.55275	38.00450	32.84530	35.70645	64.54105	
4990	2022-12-13 22:20:00	49.12230	39.84660	35.99265	37.92645	65.72715	
4999	2022-12-23 13:30:00	44.97165	35.83190	33.48865	34.23095	64.12685	

	X6	X7	X8	X9	X10	X11	X12	\
0	10.96650	145.15495	384.60800	919.95	2.05030	2.66385	3.5270	
1	10.48560	153.96970	389.11030	880.05	1.59045	1.83155	1.9555	
2	19.72395	161.72370	383.47905	712.95	1.06160	1.54015	3.5815	
3	11.15725	136.59115	379.70615	920.05	2.04615	2.74185	3.6125	
4	4.46650	149.19130	403.56375	919.95	2.18880	2.68110	3.7480	
...	
4987	7.81940	132.32080	356.39805	919.95	2.09040	2.86435	3.7385	
4988	11.28075	141.45355	322.63470	920.00	1.75145	2.33580	3.2625	

```
4989 10.52350 135.02065 357.26920 919.90 1.99050 2.47880 3.5585
4990 15.68075 138.12390 365.56105 920.00 2.05900 2.87455 2.9855
4999 22.19980 139.35930 329.82085 731.00 0.99180 1.25110 1.8820
```

```
      X13      X14      output
0    3.5270  7.69330  37.87175
1    1.9835  7.69370  35.35060
2    0.0000  7.69370  37.73035
3    3.5880  7.69310  39.72865
4    3.7475  7.69395  52.24655
...
4987  3.7295  7.69375  39.84960
4988  3.2635  7.69350  38.20440
4989  3.5740  7.69355  39.89885
4990  2.9465  7.69375  38.44980
4999  1.8485  7.69380  33.49360
```

```
[2461 rows x 16 columns]
```

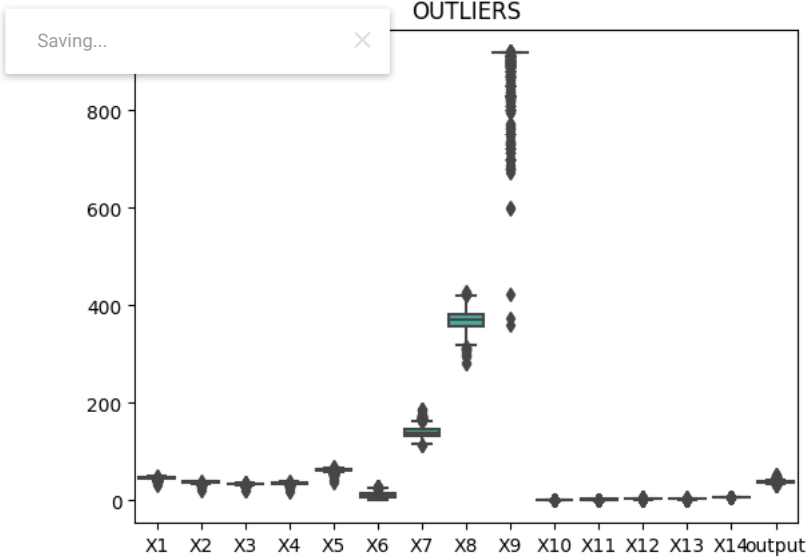
```
df.shape
```

```
(2461, 16)
```

OUTLIER DETECTION by using z score

```
numerical_columns = ['X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X7', 'X8', 'X9', 'X10', 'X11', 'X12', 'X13', 'X14', 'output']
dataf=df[numerical_columns]
sns.boxplot(data=dataf)
plt.show
plt.title('OUTLIERS')

Text(0.5, 1.0, 'OUTLIERS')
```



```
numerical_columns = ['X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X7', 'X8', 'X9', 'X10', 'X11', 'X12', 'X13', 'X14', 'output']
dataf=df[numerical_columns]
sns.scatterplot(data=dataf)
plt.show
plt.title('OUTLIERS')
```

Text(0.5, 1.0, 'OUTLIERS')



```
def find_outliers(data, threshold=3): #defining a funciton to check wheather the dataset have outlier or not

    mean = data.mean()
    std = data.std()
    z_scores = (data - mean) / std
    outliers = data[np.abs(z_scores) > threshold]

    return outliers
numerical_cols = ['CLOCK', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X7', 'X8', 'X9', 'X10', 'X11', 'X12', 'X13', 'X14', 'output']
numeric_data = df[numerical_cols]

outliers = find_outliers(numeric_data)
print(outliers)
```


	CLOCK	X1	X2	X3	X4	X5	X6	X7	X8	X9	X10	X11	X12	X13	\
0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
2	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	712.95	1.0616	1.54015	NaN	0.0	
3	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
4	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
...	
4987	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
4988	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
										NaN	NaN	NaN	NaN	NaN	
										NaN	NaN	NaN	NaN	NaN	
										NaN	NaN	NaN	NaN	NaN	

Saving...

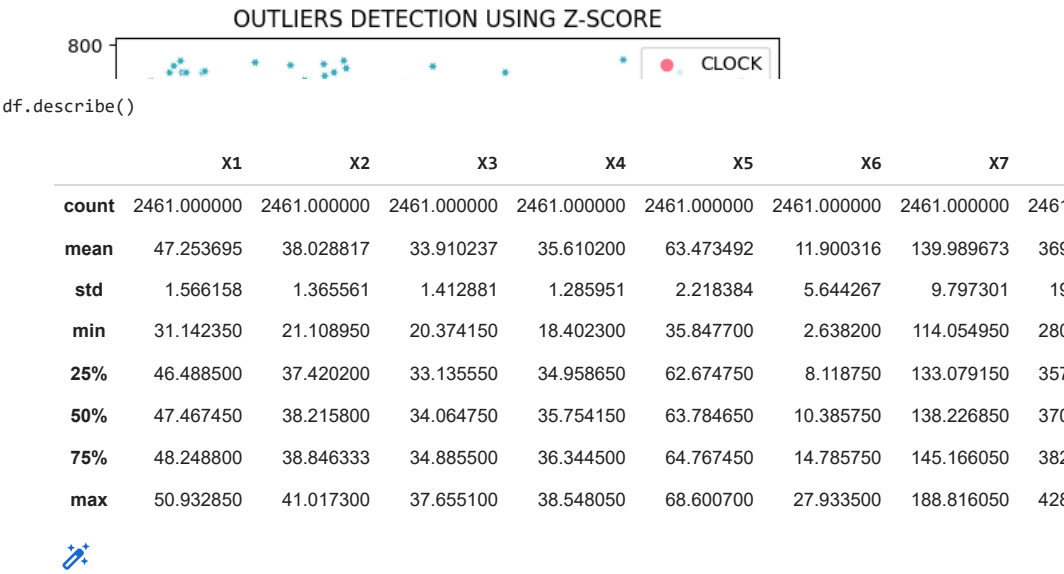
✕

	X14	output
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	52.24655
...
4987	NaN	NaN
4988	NaN	NaN
4989	NaN	NaN
4990	NaN	NaN
4999	NaN	NaN

```
[2461 rows x 16 columns]
<ipython-input-20-51f23597b19d>:3: FutureWarning: The default value of numeric_only in DataFrame.mean is deprecated. In a future ver
mean = data.mean()
<ipython-input-20-51f23597b19d>:4: FutureWarning: The default value of numeric_only in DataFrame.std is deprecated. In a future ver
std = data.std()
```




```
sns.scatterplot(data=outliers)
plt.title('OUTLIERS DETECTION USING Z-SCORE')
plt.legend()
plt.show()
```



OUTLIER REMOVAL BY IQR METHOD

```
sns.boxplot(data=df, x='X1')
```




```
df['X1'].skew()
percentile25=df['X1'].quantile(0.25)
percentile75=df['X1'].quantile(0.75)
iqr=percentile75-percentile25
upper_limit=percentile75+1.5*iqr
lower_limit=percentile25+1.5*iqr

newdf=df.loc[(df['X1']<upper_limit)&(df['X1']>lower_limit)]
print(len(df))
print(len(newdf))
```

2461
175

```
sns.boxplot(data=newdf, x='X1')
```

<Axes: xlabel='X1'>




```
df['X2'].skew()
percentile25=df['X2'].quantile(0.25)
percentile75=df['X2'].quantile(0.75)
iqr=percentile75-percentile25
upper_limit=percentile75+1.5*iqr
lower_limit=percentile25+1.5*iqr

newdf=df.loc[(df['X2']<upper_limit)&(df['X2']>lower_limit)]
print(len(df))
print(len(newdf))
```

2461
198

Saving...

<Axes: xlabel='X2'>



```
sns.boxplot(data=df, x='X3')
```


<Axes: xlabel='X3'>



```
df['X3'].skew()
percentile25=df['X3'].quantile(0.25)
percentile75=df['X3'].quantile(0.75)
iqr=percentile75-percentile25
upper_limit=percentile75+1.5*iqr
lower_limit=percentile25+1.5*iqr


newdf=df.loc[(df['X3']<upper_limit)&(df['X3']>lower_limit)]
print(len(df))
print(len(newdf))
```

2461

180


```
sns.boxplot(data=newdf, x='X3')
```

<Axes: xlabel='X3'>



```
sns.boxplot(df['X4'])
```

<Axes: >



```
df['X4'].skew()
percentile25=df['X4'].quantile(0.25)
percentile75=df['X4'].quantile(0.75)
iqr=percentile75-percentile25
```


```
upper_limit=percentile75+1.5*iqr
lower_limit=percentile25+1.5*iqr

newdf=df.loc[(df['X4']<upper_limit)&(df['X4']>lower_limit)]
print(len(df))
print(len(newdf))
```

```
2461
243
```

```
sns.boxplot(data=newdf, x='X4')
```


<Axes: xlabel='X4'>



Saving...

```
sns.boxplot(df['X5'])
```

<Axes: >




```
df['X5'].skew()
percentile25=df['X5'].quantile(0.25)
percentile75=df['X5'].quantile(0.75)
iqr=percentile75-percentile25
upper_limit=percentile75+1.5*iqr
lower_limit=percentile25+1.5*iqr

newdf=df.loc[(df['X5']<upper_limit)&(df['X5']>lower_limit)]
print(len(df))
print(len(newdf))
```

```
2461
197
```


```
sns.boxplot(data=newdf, x='X5')
```

<Axes: xlabel='X5'>



```
sns.boxplot(df['X6'])
```

<Axes: >




```
df['X6'].skew()
percentile25=df['X6'].quantile(0.25)
percentile75=df['X6'].quantile(0.75)
iqr=percentile75-percentile25
upper_limit=percentile75+1.5*iqr
lower_limit=percentile25+1.5*iqr

newdf=df.loc[(df['X6']<upper_limit)&(df['X6']>lower_limit)]
print(len(df))
print(len(newdf))
```

```
2461
227
```


```
sns.boxplot(data=newdf, x='X6')
```

<Axes: xlabel='X6'>



sns.boxplot(df['X7'])

<Axes: >




```
df['X1'].skew()
percentile25=df['X7'].quantile(0.25)
percentile75=df['X7'].quantile(0.75)
iqr=percentile75-percentile25
upper_limit=percentile75+1.5*iqr
lower_limit=percentile25+1.5*iqr

newdf=df.loc[(df['X7']<upper_limit)&(df['X7']>lower_limit)]
print(len(df))
print(len(newdf))
```

2461
266


sns.boxplot(data=newdf, x='X7')

<Axes: xlabel='X7'>



```
sns.boxplot(df['X8'])
```

<Axes: >



```
df['X8'].skew()
percentile25=df['X8'].quantile(0.25)
percentile75=df['X8'].quantile(0.75)
iqr=percentile75-percentile25
upper_limit=percentile75+1.5*iqr
lower_limit=percentile25+1.5*iqr

newdf=df.loc[(df['X8']<upper_limit)&(df['X8']>lower_limit)]
```


Saving...



2461
234

```
sns.boxplot(data=newdf, x='X8')
```

<Axes: xlabel='X8'>



```
sns.boxplot(data=df, x='X9')
```

<Axes: xlabel='X9'>



```
df['X9'].skew()
percentile25=df['X9'].quantile(0.25)
percentile75=df['X9'].quantile(0.75)
iqr=percentile75-percentile25
upper_limit=percentile75+1.5*iqr
lower_limit=percentile25+1.5*iqr


newdf=df.loc[(df['X9']<upper_limit)&(df['X9']>lower_limit)]
print(len(df))
print(len(newdf))
```

2461
189

```
sns.boxplot(data=newdf, x='X9')
```


<Axes: xlabel='X9'>

Saving... ×



```
sns.boxplot(df['X10'])
```

<Axes: >



```

df['X10'].skew()
percentile25=df['X10'].quantile(0.25)
percentile75=df['X10'].quantile(0.75)
iqr=percentile75-percentile25
upper_limit=percentile75+1.5*iqr
lower_limit=percentile25+1.5*iqr

newdf=df.loc[(df['X10']<upper_limit)&(df['X10']>lower_limit)]
print(len(df))
print(len(newdf))

```


```

2461
150

```


```
sns.boxplot(data=newdf, x='X10')
```

<Axes: xlabel='X10'>



```
sns.boxplot(df['X11'])
```

<Axes: >



```

df['X11'].skew()
percentile25=df['X11'].quantile(0.25)
percentile75=df['X11'].quantile(0.75)
iqr=percentile75-percentile25
upper_limit=percentile75+1.5*iqr
lower_limit=percentile25+1.5*iqr

newdf=df.loc[(df['X11']<upper_limit)&(df['X11']>lower_limit)]
print(len(df))
print(len(newdf))

```


```

2461
108

```


```
sns.boxplot(data=newdf, x='X11')
```

<Axes: xlabel='X11'>



```
sns.boxplot(df['X12'])
```

<Axes: >




```
df['X12'].skew()
percentile25=df['X12'].quantile(0.25)
percentile75=df['X12'].quantile(0.75)
iqr=percentile75-percentile25
upper_limit=percentile75+1.5*iqr
lower_limit=percentile25+1.5*iqr

newdf=df.loc[(df['X12']<upper_limit)&(df['X12']>lower_limit)]
print(len(df))
print(len(newdf))
```

2461
23


```
sns.boxplot(data=newdf, x='X12')
```


<Axes: xlabel='X12'>



sns.boxplot(df['X13'])

<Axes: >




```
df['X13'].skew()
percentile25=df['X13'].quantile(0.25)
percentile75=df['X13'].quantile(0.75)
iqr=percentile75-percentile25
upper_limit=percentile75+1.5*iqr
lower_limit=percentile25+1.5*iqr

newdf=df.loc[(df['X13']<upper_limit)&(df['X13']>lower_limit)]
print(len(df))
print(len(newdf))
```


```
2461
21
```

sns.boxplot(data=newdf, x='X13')

<Axes: xlabel='X13'>



```
sns.boxplot(df['X14'])
```




```
df['X14'].skew()
percentile25=df['X14'].quantile(0.25)
percentile75=df['X14'].quantile(0.75)
iqr=percentile75-percentile25
upper_limit=percentile75+1.5*iqr
lower_limit=percentile25+1.5*iqr

newdf=df.loc[(df['X14']<upper_limit)&(df['X14']>lower_limit)]
```


Saving... X

2461
0

```
sns.boxplot(data=newdf, x='X14')
```



```
sns.boxplot(df['output'])
```




```
df['output'].skew()
percentile25=df['output'].quantile(0.25)
percentile75=df['output'].quantile(0.75)
iqr=percentile75-percentile25
upper_limit=percentile75+1.5*iqr
lower_limit=percentile25+1.5*iqr

newdf=df.loc[(df['output']<upper_limit)&(df['output']>lower_limit)]
print(len(df))
print(len(newdf))
```

```
2461
137

sns.boxplot(data=newdf, x='output')
```




newdf

	CLOCK	X1	X2	X3	X4	X5	X6	X7	X8	X9
5	2022-12-26 03:00:00	46.49890	36.73950	34.69115	35.00760	60.91395	8.79055	155.91180	387.32905	919.85
32	2022-12-20 01:00:00	48.78120	39.10060	34.89470	36.61255	66.95260	9.58575	138.03350	389.17670	920.00
38	2022-12-29 07:30:00	46.57750	37.72965	33.14280	35.49450	64.09080	13.58100	134.12380	361.28640	920.00
46	2022-12-30 02:20:00	45.03310	36.36985	32.58205	34.48370	59.32215	7.38100	143.52900	336.35205	920.00
49	2022-12-19 13:50:00	47.80290	38.69705	34.91170	35.98200	65.99985	7.82405	138.09640	377.56425	919.80
...
	2022-									

```
newdf.plot.scatter(x='X1', y='output')
plt.xlabel('Independent Variable')
plt.ylabel('Dependent Variable')
```

```
plt.title('Scatter Plot: Independent vs Dependent Variable')
plt.show()
```



CHECKING STATIONARITY

```
from statsmodels.tsa.stattools import adfuller
```

```
result = adfuller(newdf['X1'])
```

Saving...

```
if p_value <= 0.05:
    print("The time series is stationary.")
```

```
ADF test p-value: 5.745986305315026e-21
The time series is stationary.
```

```
result = adfuller(newdf['X2'])
p_value = result[1]
print("ADF test p-value:", p_value)
```

```
if p_value <= 0.05:
    print("The time series is stationary.")
```

```
ADF test p-value: 1.3674325296808186e-21
The time series is stationary.
```

```
result = adfuller(newdf['X3'])
p_value = result[1]
print("ADF test p-value:", p_value)
```

```
if p_value <= 0.05:
    print("The time series is stationary.")
```

```
ADF test p-value: 1.0011844372785953e-20
The time series is stationary.
```

```
result = adfuller(newdf['X4'])
p_value = result[1]
print("ADF test p-value:", p_value)
if p_value <= 0.05:
    print("The time series is stationary.")
```

```
ADF test p-value: 1.462674513731191e-21
The time series is stationary.
```

```
result = adfuller(newdf['X5'])
p_value = result[1]
print("ADF test p-value:", p_value)
if p_value <= 0.05:
    print("The time series is stationary.")
else:
    print("The time series is notstationary.")
```

```
ADF test p-value: 5.516367337689574e-22
The time series is stationary.
```

```
result = adfuller(newdf['X6'])
p_value = result[1]
print("ADF test p-value:", p_value)
if p_value <= 0.05:
    print("The time series is stationary.")
```

```
ADF test p-value: 3.920498714625382e-21
The time series is stationary.
```

```
result = adfuller(newdf['X7'])
p_value = result[1]
print("ADF test p-value:", p_value)
if p_value <= 0.05:
    print("The time series is stationary.")
```

```
ADF test p-value: 1.465504552333475e-07
The time series is stationary.
```

```
result = adfuller(newdf['X8'])
p_value = result[1]
print("ADF test p-value:", p_value)
if p_value <= 0.05:
    print("The time series is stationary.")
```

```
ADF test p-value: 3.701978429574922e-19
The time series is stationary.
```

```
result = adfuller(newdf['X9'])
p_value = result[1]
print("The time series is stationary.")
```

```
ADF test p-value: 6.483972675406812e-21
The time series is stationary.
```

```
result = adfuller(newdf['X10'])
p_value = result[1]
print("ADF test p-value:", p_value)
if p_value <= 0.05:
    print("The time series is stationary.")
```

```
ADF test p-value: 8.499578137477644e-17
The time series is stationary.
```

```
result = adfuller(newdf['X11'])
p_value = result[1]
print("ADF test p-value:", p_value)
if p_value <= 0.05:
    print("The time series is stationary.")
```

```
ADF test p-value: 4.817530518378106e-21
The time series is stationary.
```

```
result = adfuller(newdf['X12'])
p_value = result[1]
print("ADF test p-value:", p_value)
if p_value <= 0.05:
    print("The time series is stationary.")
```

```
ADF test p-value: 6.870695129139611e-20
The time series is stationary.
```

```
result = adfuller(newdf['X13'])
p_value = result[1]
print("ADF test p-value:", p_value)
if p_value <= 0.05:
    print("The time series is stationary.")
```

```
ADF test p-value: 4.556792446926935e-23
The time series is stationary.
```

```
result = adfuller(df['X14'])
p_value = result[1]
print("ADF test p-value:", p_value)
if p_value <= 0.05:
    print("The time series is stationary.")

ADF test p-value: 0.0
The time series is stationary.

correlation=newdf.corr()
print(correlation)
```

	X1	X2	X3	X4	X5	X6	X7	\
X1	1.000000	0.514831	0.327415	0.407710	0.545590	-0.058986	-0.363575	
X2	0.514831	1.000000	0.839376	0.928894	0.805550	-0.054538	0.024117	
X3	0.327415	0.839376	1.000000	0.892398	0.579092	-0.030495	0.410912	
X4	0.407710	0.928894	0.892398	1.000000	0.760239	-0.031635	0.128197	
X5	0.545590	0.805550	0.579092	0.760239	1.000000	-0.091912	-0.205320	
X6	-0.058986	-0.054538	-0.030495	-0.031635	-0.091912	1.000000	-0.001052	
X7	-0.363575	0.024117	0.410912	0.128197	-0.205320	-0.001052	1.000000	
X8	0.194413	0.477915	0.641767	0.501176	0.281711	0.012600	0.623369	
X9	0.323272	0.183526	0.093327	0.210846	0.210467	0.124879	-0.324439	
X10	0.521576	0.405116	0.232883	0.371852	0.354414	0.078284	-0.275992	
X11	0.345737	0.265453	0.168345	0.232029	0.189636	-0.068958	-0.175923	
X12	0.216819	0.273203	0.195061	0.268187	0.099261	0.118027	-0.051588	
X13	0.160047	0.206203	0.134088	0.193876	0.134144	-0.070946	-0.016231	
X14	-0.081586	-0.139709	-0.147480	-0.116103	-0.080135	0.153524	-0.037462	
output	-0.002562	-0.068807	-0.042989	-0.057900	-0.114618	0.035996	0.064654	

	X8	X9	X10	X11	X12	X13	X14	\
X1	0.194413	0.323272	0.521576	0.345737	0.216819	0.160047	-0.081586	
X2	0.477915	0.183526	0.405116	0.265453	0.273203	0.206203	-0.139709	
X3	0.641767	0.093327	0.232883	0.168345	0.195061	0.134088	-0.147480	
X4	0.501176	0.210846	0.371852	0.232029	0.268187	0.193876	-0.116103	
X5	0.281711	0.210467	0.354414	0.189636	0.099261	0.134144	-0.080135	
X6	0.012600	0.124879	0.078284	-0.068958	0.118027	-0.070946	0.153524	
X7	0.623369	-0.324439	-0.275992	-0.175923	-0.051588	-0.016231	-0.037462	
X8	1.000000	0.607546	1.000000	0.736117	0.559524	0.724307	-0.023356	
X9	0.607546	1.000000	0.392958	0.736117	1.000000	0.598965	0.699599	
X10	0.232883	0.392958	1.000000	0.598965	0.598965	1.000000	0.486341	
X11	0.168345	0.559524	0.598965	1.000000	0.486341	1.000000	-0.139195	
X12	0.273203	0.206203	0.134088	0.193876	1.000000	0.064252	0.037144	
X13	-0.051588	-0.016231	-0.037462	-0.023356	-0.106642	-0.243951	-0.139195	
X14	-0.037462	-0.023356	-0.106642	-0.243951	-0.139195	1.000000		
output	-0.056289	-0.138181	-0.033596	0.064076	0.060252	0.064252	0.037144	


	output
X1	-0.002562
X2	-0.068807
X3	-0.042989
X4	-0.057900
X5	-0.114618
X6	0.035996
X7	0.064654
X8	-0.056289
X9	-0.138181
X10	-0.033596
X11	0.064076
X12	0.060252
X13	0.064252
X14	0.037144
output	1.000000

<ipython-input-83-987ea4cd76c3>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future ve correlation=newdf.corr()

CHECKING CORRELATION

```
plt.title('Correlation',y=1,size=16)
sns.heatmap(correlation,square=True)
```

<Axes: title={'center': 'Correlation'}>




```
x=(newdf.drop(['output','CLOCK'],axis=1))
```

```
y=(newdf['output'])
print(x.shape)
print(y.shape)
```

```
(137, 14)
(137,)
```


```
count_classes=pd.value_counts(newdf['output'],sort=True)
count_classes.plot(kind='bar',rot=0)
```

<Axes: >




```
for col in df.select_dtypes(include='number'):
    sns.regplot(x=col, y='output', data=df)
    plt.xlabel('Independent Variable')
    plt.ylabel('Dependent Variable')
    plt.title('Scatter Plot with Regression Line: Independent vs Dependent Variable')
    plt.show()
```


Scatter Plot with Regression Line: Independent vs Dependent Variable



Scatter Plot with Regression Line: Independent vs Dependent Variable




Scatter Plot with Regression Line: Independent vs Dependent Variable




Scatter Plot with Regression Line: Independent vs Dependent Variable






Scatter Plot with Regression Line: Independent vs Dependent Variable




Saving... X


Scatter Plot with Regression Line: Independent vs Dependent Variable




Scatter Plot with Regression Line: Independent vs Dependent Variable




Scatter Plot with Regression Line: Independent vs Dependent Variable



Scatter Plot with Regression Line: Independent vs Dependent Variable




Scatter Plot with Regression Line: Independent vs Dependent Variable




Scatter Plot with Regression Line: Independent vs Dependent Variable







Scatter Plot with Regression Line: Independent vs Dependent Variable




Scatter Plot with Regression Line: Independent vs Dependent Variable



Scatter Plot with Regression Line: Independent vs Dependent Variable





Scatter Plot with Regression Line: Independent vs Dependent Variable

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=42)
X_train.shape,X_test.shape
```

```
((95, 14), (42, 14))
In 45:1
```

APPLYING NORMALISATION

```
from sklearn.preprocessing import MinMaxScaler
scaler=MinMaxScaler()
scaler.fit(X_train)
X_train_scaled=scaler.transform(X_train)
X_test_scaled=scaler.transform(X_test)
```

```
X_train_scaled=pd.DataFrame(X_train_scaled,columns=X_train.columns)
X_test_scaled=pd.DataFrame(X_test_scaled,columns=X_test.columns)
```

```
np.round(X_train.describe(),1)
```

Saving...


	x1	x2	x3	x4	x5	x6	x7	x8	x9	x10	x11	x12	x13	x14
mean	46.8	37.4	33.6	35.0	61.7	12.2	141.7	371.7	913.3	1.9	2.6	3.3	3.2	7.7
std	2.3	2.7	2.4	2.6	4.9	5.6	12.0	26.0	26.3	0.3	0.3	0.6	0.7	0.0
min	35.2	21.1	20.4	18.4	35.8	3.1	120.5	284.6	762.2	1.1	1.5	1.7	0.0	7.7
25%	46.0	36.9	32.5	34.5	61.2	8.8	133.0	357.4	919.9	1.8	2.6	2.9	2.9	7.7
50%	47.3	37.9	33.9	35.6	63.1	10.5	139.6	375.8	920.0	2.0	2.7	3.5	3.4	7.7
75%	48.4	38.8	35.0	36.2	64.2	14.6	148.7	388.8	920.0	2.1	2.8	3.7	3.7	7.7
max	50.5	40.7	37.2	38.4	67.0	27.7	188.8	428.7	920.2	2.2	3.0	5.6	4.0	7.7

```
fig,(ax1, ax2)= plt.subplots(ncols=2,figsize=(12, 5))

#before scaling axi.set_title('Before Scaling')

sns.kdeplot(X_train['X1'],ax=ax1)
sns.kdeplot(X_train['X2'],ax=ax1)
sns.kdeplot(X_train['X3'],ax=ax1)
#after scaling
ax2.set_title('After Standard Scaling')
sns.kdeplot(X_train_scaled['X1'],ax=ax2)
sns.kdeplot(X_train_scaled['X2'],ax=ax2)
sns.kdeplot(X_train_scaled['X3'],ax=ax2)

plt.show()
```



ML MODEL USING LINEAR REGRESSION

```

from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
import numpy as np
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Ridge
model = LinearRegression()

param_grid = {'alpha': [0.1, 1.0, 10.0], 'max_iter': [10, 100, 1000], 'fit_intercept': [True, False]}
ridge = Ridge(param_grid)

#cross-validation
grid_search = GridSearchCV(ridge, param_grid, cv=5)
grid_search.fit(X_train_scaled, y_train)

print("hyperparameters:", grid_search.best_params_)
print("score:", grid_search.best_score_)

model = grid_search.best_estimator_
y_pred = model.predict(X_test_scaled)

mse = mean_squared_error(y_test, y_pred)

intercept = model.intercept_

print("Mean Squared Error:", mse)
print("Root Mean Squared Error:", rmse1)
print("Coefficients:", coefficients)
print("Intercept:", intercept)
print(y_pred)

hyperparameters: {'alpha': 10.0, 'fit_intercept': True, 'max_iter': 10}
score: -0.09888222757513718
Mean Squared Error: 0.39221779052798167
Root Mean Squared Error: 0.6262729361292739
Coefficients: [-0.01153665 -0.02485005  0.00069192 -0.01478967 -0.05382455  0.04310128
  0.10618496 -0.02874678 -0.11032604 -0.05280059 -0.0331805  -0.00433651
  0.01812458  0.0482288 ]
Intercept: 41.196426214273984
[41.01036894 41.02120848 41.04439714 41.02290897 41.1805145  41.04303165
 40.98302172 40.96117219 41.1582513  41.0053333  41.02816517 41.01113833
 40.99259039 41.2325796  41.06444147 41.02082327 40.99231579 40.9910005
 41.06584699 40.99168217 41.04286046 40.95769599 41.03297301 40.99586992
 40.99512223 41.0272108  40.98747652 41.02186067 41.00108212 40.99164623
 41.00984398 40.99552306 41.06183716 41.02138876 41.00831152 41.01250068
 40.95507321 41.02847313 41.01660515 40.97256688 41.01351563 40.96107077]

import matplotlib.pyplot as plt


plt.scatter(range(len(y_test)), y_test)
plt.scatter(range(len(y_pred)), y_pred)
plt.legend(['ACTUAL', 'PREDICTED'])
plt.show()

```



```
residuals = y_test - y_pred
```

```
plt.scatter(y_pred, residuals)
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.axhline(y=0, color='r', linestyle='--')
plt.title('Residual Plot')
plt.show()
```



ML MODEL USING XGBOOST

```
import xgboost as xgb
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import numpy as np
param_grid = {
    'max_depth': [3, 5, 7],
    'learning_rate': [0.1, 0.01, 0.001],
    'n_estimators': list(range(100, 1000, 100))
}
model = xgb.XGBRegressor(objective='reg:squarederror', max_depth=3, learning_rate=0.1, n_estimators=100)

grid_search = GridSearchCV(model, param_grid, scoring='neg_mean_squared_error', cv=5)
grid_search.fit(X_train_scaled, y_train)

print("Hyperparameters:", grid_search.best_params_)
print("Best score:", -grid_search.best_score_)
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test_scaled)

mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
rmse2 = np.sqrt(mse)


print("Mean Squared Error:", mse)
print("Mean Absolute Error:", mae)
print("Root Mean Squared Error:", rmse2)

r2 = r2_score(y_test, y_pred)
print("R-squared Score:", r2)
```


```
Hyperparameters: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 600}
Best score: 0.41745777882228535
Mean Squared Error: 0.4085904635994367
Mean Absolute Error: 0.47753576543898746
```

Root Mean Squared Error: 0.6392108131121036
 R-squared Score: -0.06538119922619368

```
plt.figure(figsize=(10, 6))
plt.scatter(range(len(y_test)), y_test, color='b', label='Actual')
plt.scatter(range(len(y_pred)), y_pred, color='r', label='Predicted')
plt.xlabel('Index')
plt.ylabel('Value')
plt.title('Actual vs. Predicted Values')
plt.legend()
plt.show()
```



```
residuals = y_test - y_pred
plt.scatter(y_pred, residuals)
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.axhline(y=0, color='r', linestyle='--')
plt.title('Residual Plot')
plt.show()
```



ML MODEL USING SUPPORT VECTOR REGRESSION

```
from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
```

```

from sklearn.model_selection import GridSearchCV

model = SVR(kernel='rbf')
param_grid = {
    'C': [0.1, 1, 10],
    'gamma': [0.01, 0.1, 1],
    'epsilon': [0.01, 0.1, 1],
    'kernel': ['rbf', 'linear', 'poly']
}
# Perform grid search with cross-validation
grid_search = GridSearchCV(model, param_grid, scoring='neg_mean_squared_error', cv=5)
grid_search.fit(X_train_scaled, y_train)

best_model = grid_search.best_estimator_

y_pred = best_model.predict(X_test_scaled)
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse3 = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print("Mean Absolute Error:", mae)
print("Hyperparameters:", grid_search.best_params_)
print("Best score:", -grid_search.best_score_)
print("Root Mean Squared Error:", rmse3)
print("Mean Squared Error:", mse)
print("R-squared:", r2)

Mean Absolute Error: 0.43944475567837415
Hyperparameters: {'C': 1, 'epsilon': 0.01, 'gamma': 1, 'kernel': 'rbf'}
Best score: 0.3817036498095305
Root Mean Squared Error: 0.5916615002961743
Mean Squared Error: 0.3500633309327198
R-squared: 0.08722565860019482


```

Saving...

```

plt.scatter(range(len(y_test)), y_test, color='b', label='Actual')
plt.scatter(range(len(y_pred)), y_pred, color='r', label='Predicted')
plt.xlabel('Index')
plt.ylabel('Value')
plt.title('Actual vs. Predicted Values')
plt.legend()
plt.show()

```




```

residuals = y_test - y_pred
plt.scatter(y_pred, residuals)
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.axhline(y=0, color='r', linestyle='--')

```



```
plt.title('Residual Plot')
plt.show()
```



ML MODEL USING RANDOM FOREST REGRESSOR

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split, GridSearchCV

Saving... (1000, 100)),
'max_depth': [None, 5, 10],
'min_samples_split': [2, 5, 10],
'min_samples_leaf': [1, 2, 4]
}
model = RandomForestRegressor(random_state=42)
# Perform grid search with cross-validation
grid_search = GridSearchCV(model, param_grid, scoring='neg_mean_squared_error', cv=5)
grid_search.fit(X_train_scaled, y_train)


best_model = grid_search.best_estimator_

y_pred = best_model.predict(X_test_scaled)
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse4 = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)


print("Mean Absolute Error:", mae)
print("Hyperparameters:", grid_search.best_params_)
print("Best score:", -grid_search.best_score_)
print("Root Mean Squared Error:", rmse4)
print("Mean Squared Error:", mse)
print("R-squared:", r2)

Mean Absolute Error: 0.4640138351644905
Hyperparameters: {'max_depth': 5, 'min_samples_leaf': 4, 'min_samples_split': 10, 'n_estimators': 700}
Best score: 0.37442932730958683
Root Mean Squared Error: 0.5846155985097846
Mean Squared Error: 0.34177539802095364
R-squared: 0.10883606973622173

plt.figure(figsize=(10, 6))
plt.scatter(range(len(y_test)), y_test, color='b', label='Actual')
plt.scatter(range(len(y_pred)), y_pred, color='r', label='Predicted')
plt.xlabel('Index')
plt.ylabel('Value')
plt.title('Actual vs. Predicted Values')
plt.legend()
plt.show()
```



```
residuals = y_test - y_pred
plt.scatter(y_pred, residuals)
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.axhline(y=0, color='r')
plt.title('Residual Plot')
plt.show()
```



MODEL ACCURACY COMPARISON

```
models = ['Linear Regression', 'XGBoost', 'SVR', 'Random Forest']
rmse = [rmse1, rmse2, rmse3, rmse4]
sns.set(style="whitegrid")
plt.figure(figsize=(10, 6))
sns.barplot(x=models, y=rmse)
plt.title('Accuracy of Different Models')
plt.xlabel('Models')
plt.ylabel('RMSE')
plt.show()
```



From the following bar plots we got to know that XGBOOST gives the most accurate prediction